
title: " Coordinate Descent for Lasso"

author: "Harini"

date: "15 September 2020"

output: Doc_document

Coordinate Descent for Lasso

First, prepare the Boston Housing Data. Check [Stat542_ Variable Selection.html] on relevant background information.

```
library(MASS)
```

```
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loading required package: foreach
```

```
## Loaded glmnet 2.0-16
```

```
myData = Boston
```

```
names(myData)[14] = "Y"
```

```
iLog = c(1, 3, 5, 6, 8, 9, 10, 14);
```

```
myData[, iLog] = log(myData[, iLog]);
```

```
myData[, 2] = myData[, 2] / 10;
```

```
myData[, 7] = myData[, 7]^2.5 / 10^4
```

```
myData[, 11] = exp(0.4 * myData[, 11]) / 1000;
```

```
myData[, 12] = myData[, 12] / 100;
```

```
myData[, 13] = sqrt(myData[, 13]);
```

```
X = as.matrix(myData[, -14]) y
```

```
= myData$Y
```

```
lam.seq = c(0.30, 0.2, 0.1, 0.05, 0.02, 0.005)
```

Next write my own function to implement CD, which should output estimated Lasso coefficients similar to the one output by R.

```
MyLasso = function(X, y, lam.seq, maxit, standardize){
```

```
  # X: n-by-p design matrix without the intercept
```

```
  # y: n-by-1 response vector
```

```
  # lam.seq: sequence of lambda values
```

```
  # maxit: number of updates for each lambda
```

```
  # standardize: if True, center and scale X and y.
```

```
  n = length(y)
```

```
  p = dim(X)[2]
```

```
  nlam = length(lam.seq)
```

```
  if(standardize==TRUE){
```

```
    Xmean = apply(X,2,mean)
```

```
    Xsd = apply(X,2,sd)*sqrt((n-1)/n)
```

```
    X = t((t(X)-Xmean)/Xsd)
```

```
    ymean = mean(y)
```

```

ysd = sd(y)*sqrt((n-1)/n)
y = (y-ymean)/ysd
# Center and scale X and y
# Record the corresponding means and scales
}

# Initilize coef vector b and residual vector r
b = rep(0, p)
r = y
B = matrix(0,nlam,p+1)

one_step_lasso = function(r, x, lam){
  xx = sum(x^2)
  xr = sum(r*x)
  b = (abs(xr) - lam/2)/xx
  b = sign(xr)*ifelse(b>0, b, 0)
  return(b)
}

# Triple nested loop
for(m in 1:nlam){
  lam = (2*n)*lam.seq[m]/ysd # assign lambda value
  for(step in 1:maxit){
    for(j in 1:p){
      r = r + (X[,j]*b[j])
      b[j] = one_step_lasso(r, X[, j], lam)
      r = r - X[, j] * b[j]
    }
  }
  B[m, -1] = b
}

if(standardize==TRUE){
  for(m in 1:nlam){
    B[m, -1] = B[m, -1]/Xsd*ysd
    B[m, 1] = ymean - sum(B[m, -1]*Xmean)
  }
  # scale back the coefficients and update the intercepts B[, 1]
}

return(t(B))
}

```

Check the accuracy of my algorithm against the output from glmnet.

```

lam.seq = c(0.30, 0.2, 0.1, 0.05, 0.02, 0.005)
lasso.fit = glmnet(X, y, alpha = 1, lambda = lam.seq, standardize = TRUE)
coef(lasso.fit)

```

```

## 14 x 6 sparse Matrix of class "dgCMatrix"
##              s0              s1              s2              s3              s4
## (Intercept)  3.16239335  3.5089461  3.855935763  3.778455800  3.542129253

```

```
## crim      .      .      .      .      .
## zn        .      .      .      .      .
## indus     .      .      .      .      .
## chas      .      .      .      .      0.066692255
## nox       .      .      .      .      .
## rm        .      .      .      0.240416063  0.417249062
## age       .      .      .      .      .
## dis       .      .      .      .      -0.004681106
## rad       .      .      .      .      .
## tax       .      .      .      -0.055308804 -0.080579126
## ptratio   .      .      -0.004310305 -0.023647910 -0.033626958
## black     .      .      .      0.008515144  0.031234249
## lstat     -0.03741741 -0.1388176 -0.237634045 -0.244558377 -0.243489936
##          s5
## (Intercept) 3.925458057
## crim      .
## zn        .
## indus     -0.005398865
## chas      0.099119798
## nox       -0.125162756
## rm        0.459272519
## age       .
## dis       -0.120437573
## rad       0.014980251
## tax       -0.149039129
## ptratio   -0.038211176
## black     0.043402206
## lstat     -0.256708941
```

```
myout = MyLasso(X, y, lam.seq, maxit = 50, standardize = TRUE)
rownames(myout) = c("Intercept", colnames(X))
myout
```

```
##          [,1]      [,2]      [,3]      [,4]      [,5]
## Intercept 3.16239335 3.5089461 3.855935052 3.777919609 3.541573070
## crim      0.00000000 0.0000000 0.000000000 0.000000000 0.000000000
## zn        0.00000000 0.0000000 0.000000000 0.000000000 0.000000000
## indus     0.00000000 0.0000000 0.000000000 0.000000000 0.000000000
## chas      0.00000000 0.0000000 0.000000000 0.000000000 0.066670285
## nox       0.00000000 0.0000000 0.000000000 0.000000000 0.000000000
## rm        0.00000000 0.0000000 0.000000000 0.240724202 0.417613831
## age       0.00000000 0.0000000 0.000000000 0.000000000 0.000000000
## dis       0.00000000 0.0000000 0.000000000 0.000000000 -0.004686483
## rad       0.00000000 0.0000000 0.000000000 0.000000000 0.000000000
## tax       0.00000000 0.0000000 0.000000000 -0.055335623 -0.080634128
## ptratio   0.00000000 0.0000000 -0.004303294 -0.023646161 -0.033638505
## black     0.00000000 0.0000000 0.000000000 0.008522442 0.031255297
## lstat     -0.03741741 -0.1388176 -0.237638248 -0.244528823 -0.243439646
##          [,6]
## Intercept 3.927045881
## crim      0.000000000
## zn        0.000000000
## indus     -0.005275075
## chas      0.099049491
## nox       -0.123495061
```

```
## rm      0.459679619
## age     0.000000000
## dis     -0.119808332
## rad     0.015118868
## tax     -0.149512107
## ptratio -0.038234008
## black   0.043452192
## lstat   -0.256682319
```

The maximum difference between the two coefficient matrices is less than 0.005.

```
max(abs(coef(lasso.fit) - myout))
```

```
## [1] 0.001667695
```